

# S\_P\_Timeseries

Vijay Mudivedu

6/6/2018

## Time Series Analysis of S&P 500

### Business Understanding

Business objective:

Please apply any machine learning algorithm you are comfortable with for predicting this time series

The results would be measured on:

1. Accuracy - How good are predictions
2. Visualization - How well are you able to convey your idea graphically
3. Code Cleanliness - How well have you documented your code in an easy language to understand. No need for excess code

Read Data from the CSV

```
## [1] 643 12

##  symbol  date_txn  open    low    high close_price    volume
lead_1
## 1    SPY  11/10/2015  207.51  207.19  208.60      208.55  71844000
207.67
## 2    SPY  11/11/2015  208.88  207.66  208.94      207.67  67251000
204.84
## 3    SPY  11/12/2015  206.50  204.82  207.06      204.84  118209400
202.54
## 4    SPY  11/13/2015  204.35  202.44  204.67      202.54  145494400
205.62
## 5    SPY  11/16/2015  202.32  202.18  205.69      205.62  112996000
205.47
## 6    SPY  11/17/2015  205.99  204.88  207.04      205.47  113429400
208.73
##  lead_5 lead_10      name class_type_of
## 1  205.47  209.35 SPDR  S&P500      S_P_500
## 2  208.73  209.32 SPDR  S&P500      S_P_500
## 3  208.55  209.56 SPDR  S&P500      S_P_500
## 4  209.31  208.69 SPDR  S&P500      S_P_500
## 5  209.07  210.68 SPDR  S&P500      S_P_500
## 6  209.35  208.53 SPDR  S&P500      S_P_500
```

```
## 'data.frame':    643 obs. of  12 variables:
## $ symbol      : Factor w/ 1 level "SPY": 1 1 1 1 1 1 1 1 1 1 ...
## $ date_txn    : Factor w/ 643 levels "1/10/2017","1/10/2018",...:
106 109 111 112 118 121 124 126 129 135 ...
## $ open        : num  208 209 206 204 202 ...
## $ low         : num  207 208 205 202 202 ...
## $ high        : num  209 209 207 205 206 ...
## $ close_price  : num  209 208 205 203 206 ...
## $ volume       : int  71844000 67251000 118209400 145494400
112996000 113429400 113064100 81363500 89556300 63829000 ...
## $ lead_1      : num  208 205 203 206 205 ...
## $ lead_5      : num  205 209 209 209 209 ...
## $ lead_10     : num  209 209 210 209 211 ...
## $ name        : Factor w/ 1 level "SPDR S&P500": 1 1 1 1 1 1 1 1 1
1 ...
## $ class_type_of: Factor w/ 1 level "S_P_500": 1 1 1 1 1 1 1 1 1 1
...
```

## Data Understanding

### Data Preparation

- Remove Unwanted Columns

Remove the redundant columns symbol, name, class\_type\_of "SPY", "SPDR S&P500", "S\_P\_500"

```
##   date_txn  open  low  high close_price  volume lead_1
lead_5
## 1 11/10/2015 207.51 207.19 208.60      208.55  71844000 207.67
205.47
## 2 11/11/2015 208.88 207.66 208.94      207.67  67251000 204.84
208.73
## 3 11/12/2015 206.50 204.82 207.06      204.84 118209400 202.54
208.55
## 4 11/13/2015 204.35 202.44 204.67      202.54 145494400 205.62
209.31
## 5 11/16/2015 202.32 202.18 205.69      205.62 112996000 205.47
209.07
## 6 11/17/2015 205.99 204.88 207.04      205.47 113429400 208.73
209.35
##   lead_10
## 1   209.35
## 2   209.32
## 3   209.56
## 4   208.69
## 5   210.68
## 6   208.53
```

- Check for NAs

```
##      date_txn      open      low      high close_price
volume
##          0          0          0          0          0
0
##      lead_1      lead_5      lead_10
##          2          6          11
```

*Comments:* NAs found in Lead\_1, Lead\_2 and Lead\_10

- Query to find NAs

```
## [1] 643 9
```

*Comments:* NAs were identified in lead\_1, lead\_5, lead\_10

- NAs in Lead\_1, Lead\_5, Lead\_10 are:

```
## $lead_1
## [1] 590 643
##
## $lead_5
## [1] 586 639 640 641 642 643
##
## $lead_10
## [1] 581 634 635 636 637 638 639 640 641 642 643
```

- Remove NAs from lead\_10

```
## [1] 632 9
```

- NAs in lead\_1 and lead\_5 are:

```
## $lead_1
## [1] 589
##
## $lead_5
## [1] 585
##
## $lead_10
## integer(0)
```

- Removing the NAs from lead\_1 remaining the rows and columns are

```
## [1] 631 9
```

- NAs in Lead 5 are:


```
## $lead_1
## integer(0)
##
## $lead_5
## [1] 585
##
## $lead_10
## integer(0)
```

- Removing the NAs from lead\_5, remaining the rows and columns are:

```
## [1] 630 9
```

```
##      date_txn   open    low   high close_price    volume lead_1
lead_5
## 1 11/10/2015 207.51 207.19 208.60      208.55 71844000 207.67
205.47
## 2 11/11/2015 208.88 207.66 208.94      207.67 67251000 204.84
208.73
## 3 11/12/2015 206.50 204.82 207.06      204.84 118209400 202.54
208.55
## 4 11/13/2015 204.35 202.44 204.67      202.54 145494400 205.62
209.31
## 5 11/16/2015 202.32 202.18 205.69      205.62 112996000 205.47
209.07
## 6 11/17/2015 205.99 204.88 207.04      205.47 113429400 208.73
209.35
##      lead_10
## 1      209.35
## 2      209.32
## 3      209.56
## 4      208.69
## 5      210.68
## 6      208.53
```

- check for outliers in close\_price



*Comments:* No outliers found

- Convert the date to R date format

```
## [1] "2015-11-10" "2015-11-11" "2015-11-12" "2015-11-13" "2015-11-16"
## [6] "2015-11-17"
```

- Check for uniuiques and duplicates

```
## [1] 0
```

*Comments:* No duplicates found

- Check for invalid valid data - checking for minimum dips

##	open	low	high	close_price	volume
lead_1					
##	382	382	382	64	484
63					
##	lead_5	lead_10			
##	377	54			

*Comments :* There is bad data in variables open, low, high, lead\_1, thus considering the “close\_price”

## Exploratory Data Analysis

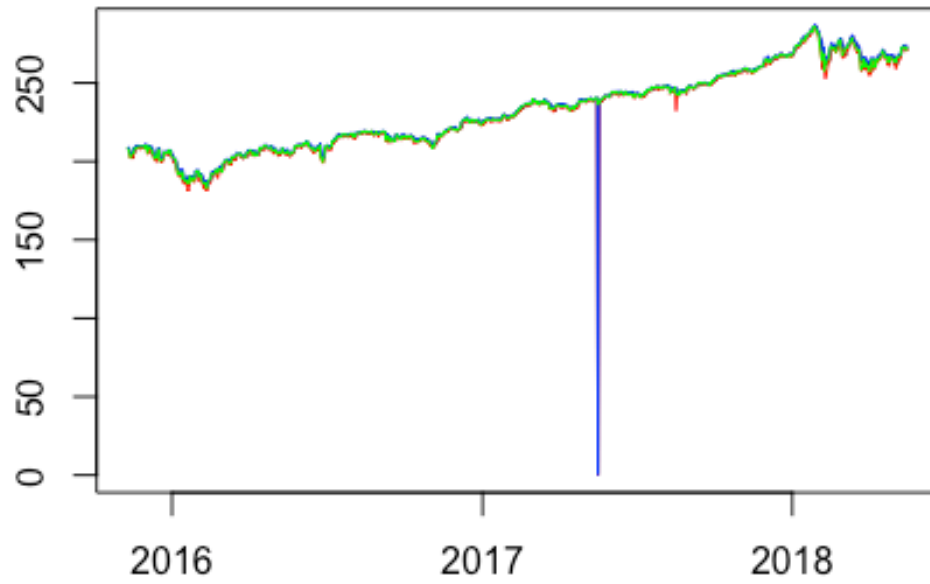
- Plotting the datasets

```
##    date_txn    open      low      high close_price
volume
##          0          0          0          0          0
0
##    lead_1    lead_5    lead_10
##          0          0          0
```



*Comments:* Thus considering the close\_price as the lead\_prices do not convey much information


## Open, low, high Price



- Univariate analysis

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

## Distribution of Close\_price Data



Comments: Most of the Pricing data is concentrated around 200 - 240

### Time Series Object: Creating a time series object of "sp\_historical\_cleaned"

```
##          open    low   high close_price    volume lead_1 lead_5
## 2015-11-10 207.51 207.19 208.60      208.55  71844000 207.67 205.47
## 2015-11-11 208.88 207.66 208.94      207.67  67251000 204.84 208.73
## 2015-11-12 206.50 204.82 207.06      204.84 118209400 202.54 208.55
## 2015-11-13 204.35 202.44 204.67      202.54 145494400 205.62 209.31
## 2015-11-16 202.32 202.18 205.69      205.62 112996000 205.47 209.07
## 2015-11-17 205.99 204.88 207.04      205.47 113429400 208.73 209.35
##          lead_10
## 2015-11-10  209.35
## 2015-11-11  209.32
## 2015-11-12  209.56
## 2015-11-13  208.69
## 2015-11-16  210.68
## 2015-11-17  208.53
```





```
##      date_txn  open    low   high close_price    volume lead_1
lead_5
## 604 2018-04-06 263.42 258.00 265.11      259.72 182029210 261.00
265.15
## 605 2018-04-09 261.37 259.94 264.84      261.00 104745500 265.15
267.33
## 606 2018-04-10 264.27 262.98 266.04      265.15 104375800 263.76
270.19
## 607 2018-04-11 263.47 263.39 265.64      263.76  90886300 265.93
270.39
## 608 2018-04-12 265.26 265.06 267.00      265.93  68138500 265.15
268.89
## 609 2018-04-13 267.41 264.01 267.54      265.15  84647281 267.33
266.61
##      lead_10
## 604 266.61
## 605 266.57
## 606 262.98
## 607 263.63
## 608 266.31
## 609 266.56

## [1] 600    9

## [1] 30    9
```

## Smoothing

### Smoothing using Classical Decomposition

- window width

```
## [1] 9
```

- Printing the smoothed moving average:

```
## Time Series:
## Start = 1
## End = 20
## Frequency = 1
## [1]      NA      NA      NA      NA      NA      NA      NA
## [8]      NA      NA 207.8979 207.8137 207.6911 207.7453 207.7058
## [15] 207.5626 207.5395 207.5026 207.3084 206.8195 206.4300
```

- Replacing the induced NAs as result of smoothing

```
## [1] -0.08421053
```

```
## [1] 0.2178947
```

- Replacing the lagging NAs are a result of windowing of moving average

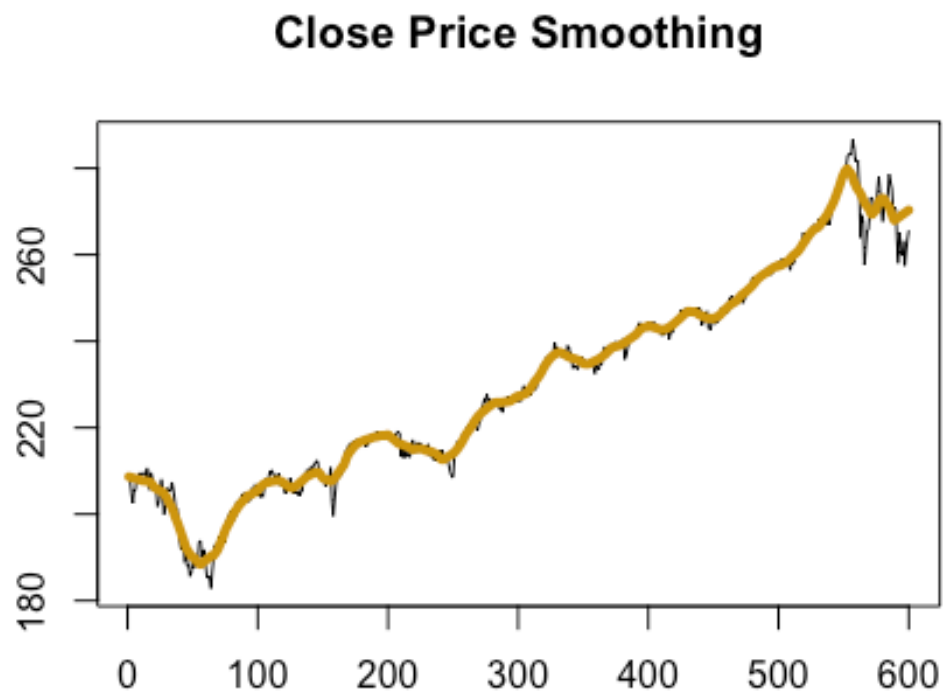
```
## Time Series:
## Start = 1
## End = 20
```

```
## Frequency = 1
## [1] 208.6558 208.5716 208.4874 208.4032 208.3189 208.2347 208.1505
## [8] 208.0663 207.9821 207.8979 207.8137 207.6911 207.7453 207.7058
## [15] 207.5626 207.5395 207.5026 207.3084 206.8195 206.4300
```

- Replacing the smoothed leading NAs


```
## Time Series:
## Start = 581
## End = 600
## Frequency = 1
## [1] 272.9321 272.4353 271.7326 271.4726 270.9568 270.1737 269.3974
## [8] 268.5047 267.9321 268.1500 268.3679 268.5858 268.8037 269.0216
## [15] 269.2395 269.4574 269.6753 269.8932 270.1111 270.3289
```

- Plotting the Smoothed Close Price



## Smoothing using Holts method

### Smoothing - Holts Method



*Comments:* Clearly, from Holts Method best smoothing happens when alpha is  $\sim 0.1$

- Holts smoothed series Plot

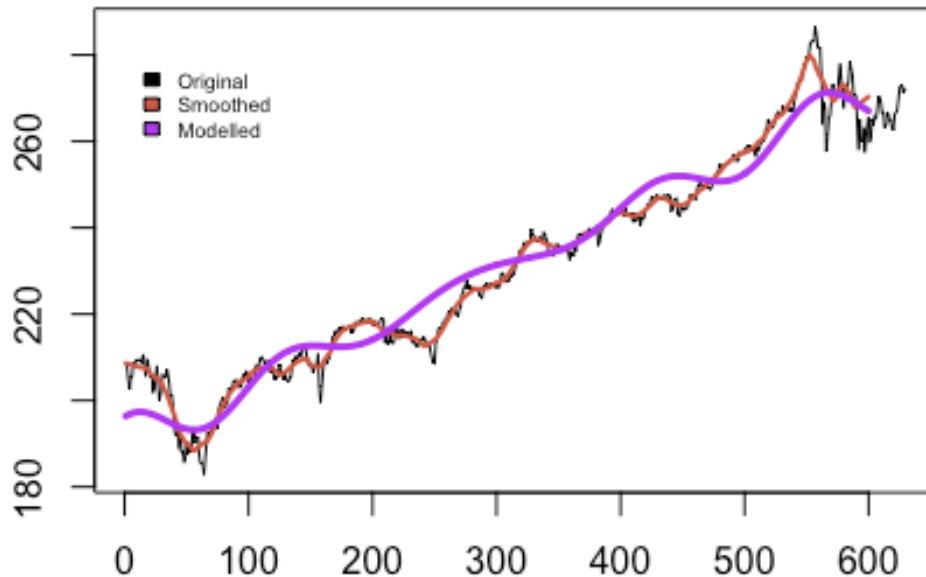


```

## (Intercept) 230.23069
0.19667
## cos(0.05 * months_smoothed) -0.02158
0.27776
## poly(months_smoothed, 1) 567.96883
4.87636
## sin(0.05 * months_smoothed) 1.85678
0.27629
## cos(0.05 * months_smoothed):poly(months_smoothed, 1) -88.38120
6.92574
## poly(months_smoothed, 1):sin(0.05 * months_smoothed) 6.63155
6.78174
## t value
Pr(>|t|)
## (Intercept) 1170.650 < 2e-
16 ***
## cos(0.05 * months_smoothed) -0.078
0.938
## poly(months_smoothed, 1) 116.474 < 2e-
16 ***
## sin(0.05 * months_smoothed) 6.720 4.25e-
11 ***
## cos(0.05 * months_smoothed):poly(months_smoothed, 1) -12.761 < 2e-
16 ***
## poly(months_smoothed, 1):sin(0.05 * months_smoothed) 0.978
0.329
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.77 on 594 degrees of freedom
## Multiple R-squared:  0.961, Adjusted R-squared:  0.9606
## F-statistic: 2925 on 5 and 594 DF, p-value: < 2.2e-16

```

## SPY Close Price



*Comments:* With adjusted R-Squared with accuracy of 96.2% is the best fit curve and it has also generated the least MAPE error of 1%.

- Stationarity tests on the residual time series:

	1	2	3	4	5	6
##	12.134755	11.095516	8.121673	5.693146	8.659823	8.411551

- adf test and kpss test for stationarity:

```
## Warning in adf.test(x = resi_close_price, alternative =
"stationary"): p-
## value smaller than printed p-value

##
## Augmented Dickey-Fuller Test
##
## data: resi_close_price
## Dickey-Fuller = -4.0685, Lag order = 8, p-value = 0.01
## alternative hypothesis: stationary

##
## KPSS Test for Level Stationarity
##
## data: resi_close_price
```


```
## KPSS Level = 0.6319, Truncation lag parameter = 5, p-value =  
## 0.01974
```

*Comments:* From these tests it can be inferred that there is enough evidence to prove that the “resi\_close\_price” is Stationary.


- ACF and PACF Plots are:



**Series resi\_close\_price**



### Series resi\_close\_price



```
## Series: resi_close_price
## ARIMA(0,0,0) with non-zero mean
##
## Coefficients:
##      mean
##    -0.0961
## s.e.   0.2227
##
## sigma^2 estimated as 29.81:  log likelihood=-1869.27
## AIC=3742.55   AICc=3742.57   BIC=3751.34
```

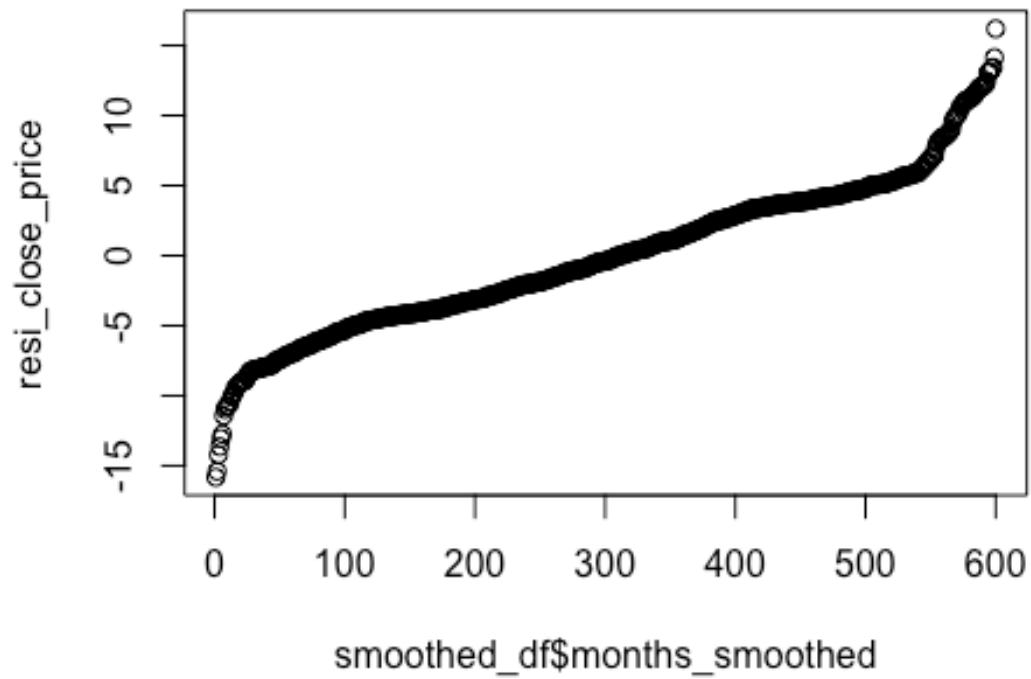
*Comments:* This has a very small sigma square, with a very high log likelihood. In addition to this this is AR(0) and MA(0) time series

- checking the noise and stationarity of the time series using the box-Ljung test

```
##
## Box-Ljung test
##
## data:  resi_close_price
## X-squared = 535.84, df = 1, p-value < 2.2e-16
```

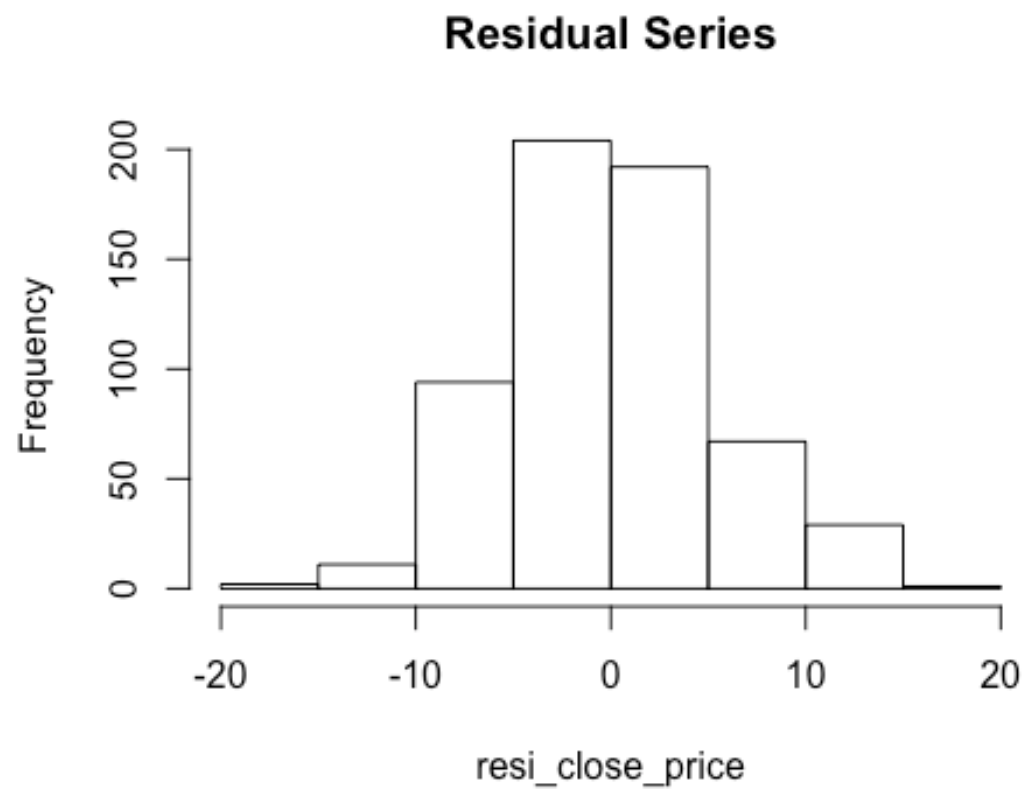
*Comments:* thus the p-value of the time-series is very low making it a good fit to call it a Strictly Stationary time series.

- QQPlot of the residual timeseries




*Comments:* QQ plot suggests that Close\_price time-series is stationary

- Plotting the histogram of the residual time-series represents a Gaussian Curve



- Checking the properties of Residual time-series by inspection

## Residual Time series



*Comments:* Thus, it is also clear from the plots ARIMA, ADF test, KPSS test the that the time series is stationary

### Model Evaluation

- Predicting the values fitted model from the vaddilation dataset

```
## months_smoothed
## 1                601
## 2                602
## 3                603
## 4                604
## 5                605
## 6                606
```

- MAPE Error - Accuracy Calculation

```
## [1] 1.026822
```

*Comments:* Thus with a MAPE error of just 1.0268, model is apparently a very good fit.

- Combining the predicted values train and test data
- Plottig the actual vs the predicted data, using classical decomposition




## PREDICTION ANALYSIS USING AUTO.ARIMA

### Modelling auto arima

```
## Series: sp_train$close_price
## ARIMA(4,1,0)
##
## Coefficients:
##          ar1      ar2      ar3      ar4
##      -0.0655  -0.0858   0.0847   0.0362
## s.e.   0.0409   0.0410   0.0411   0.0416
##
## sigma^2 estimated as 3.148:  log likelihood=-1191.36
## AIC=2392.73  AICc=2392.83  BIC=2414.7
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE
MASE
## Training set 0.09879535 1.766713 1.188871 0.03878482 0.5244785
0.999416
##              ACF1
## Training set -0.002495871
```

- Predicting the Validation dataset

## SPy Close Price



- MAPE of Close Price Using Auto Arima.

```
## [1] 1.079454
```

- Combining the predicted training and testing to plot
- Combined plot of Arima



### Conclusion

- Thus the Accuracy result using the MAPE using Classical Decomposition is 1.03 which is less than Auto.Arima Modelling 1.08 and Classical Decomposition has had higher accuracy\*